



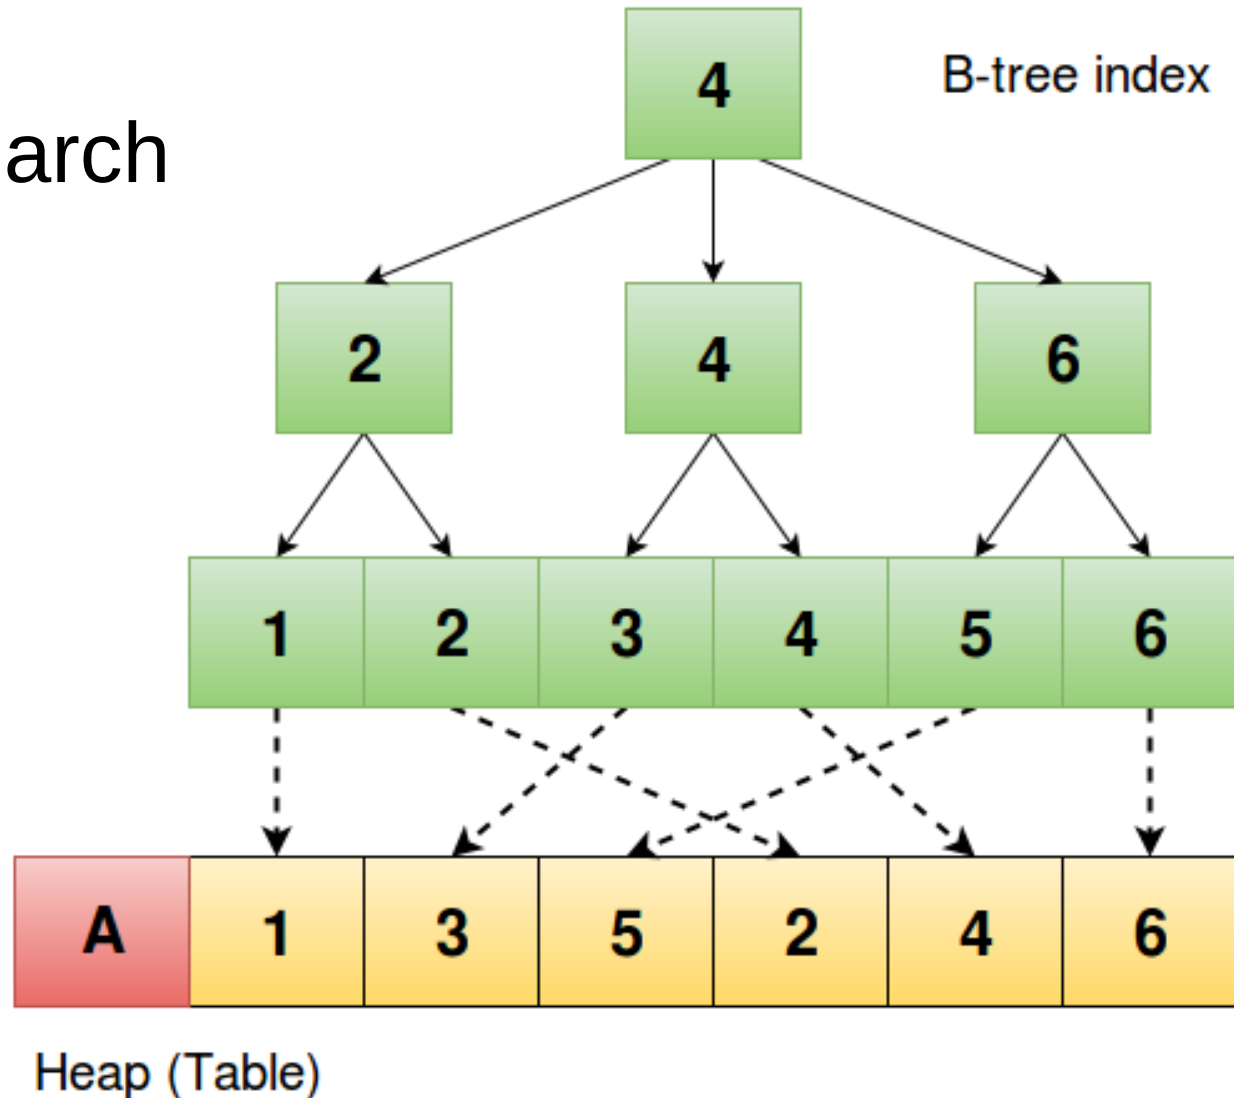
# Indexes don't mean slow inserts

Anastasia Lubennikova

1. Why do we need it?
2. Write-optimisation techniques
3. PostgreSQL specific
4. Advanced PostgreSQL indexes
5. Future of indexing in PostgreSQL


# PostgreSQL indexes

- Speed up search
- Primary key
- Constraints
  
- Secondary indexes




# Index maintenance overhead

- Index size
- INSERT slowdown
- Random I/O
- Index becomes fragmented
- More indexes - more overhead



# Do we need write-optimisation?


- Heavy write load
- MVCC  
update = insert



UPDATE mytable SET a = a + 1;

# Do we need write-optimisation?

- 1Gb table
- Update all values:  
Without index ~ 200s  
With index ~ 600s



# DBMS trade-offs



# DBMS trade-offs

- CAP-theorem
- ACID vs BASE
- Lower hardware cost vs Better productivity
- Read speed vs Write speed
- Productivity vs Fault-tolerance



# Write-optimisation

- Writes are faster
- Reads are good
- Storage is fault-tolerant

# Insert buffer

- Accumulate data. Sort. Insert at a time.

+ Avoids random I/O

- Seqscan buffer

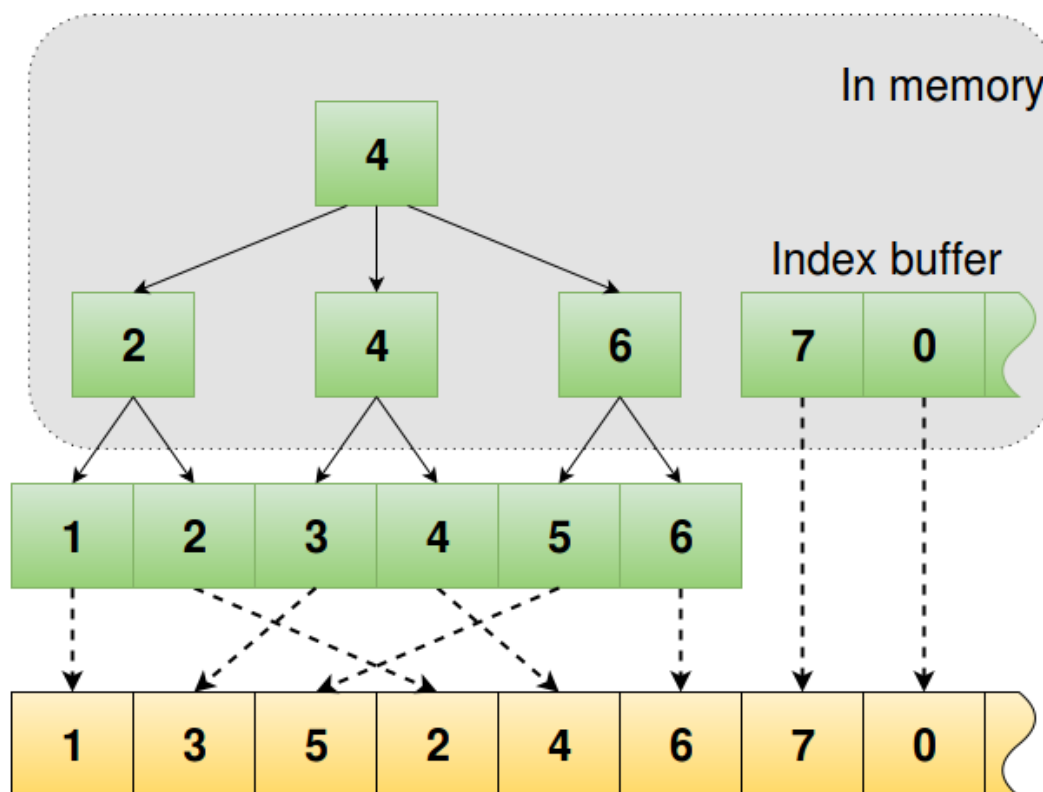
- Possible data loss

- Merge time

- Avoids hidden scans


- only non-unique

- MySQL InnoDB Change Buffer



# Cache-oblivious data structures


- Approximately optimal for any hardware
- Divide & Conquer



# LSM trees


- Cascade of B-trees
- First tree is in memory

- LevelDB
- BigTable
- Cassandra
- Hbase
- SophiaDB
- other NoSQL DBs




# From LSM to COLA

- Search optimisation for LSM
- Leaf levels are linked by lookahead pointers
- Introduced in 2007 by founders of Tokutek




# Cache-oblivious lookahead arrays

- Drop internal trees nodes
- Bound the scan area with lookahead pointers



# COLA: theory and practice

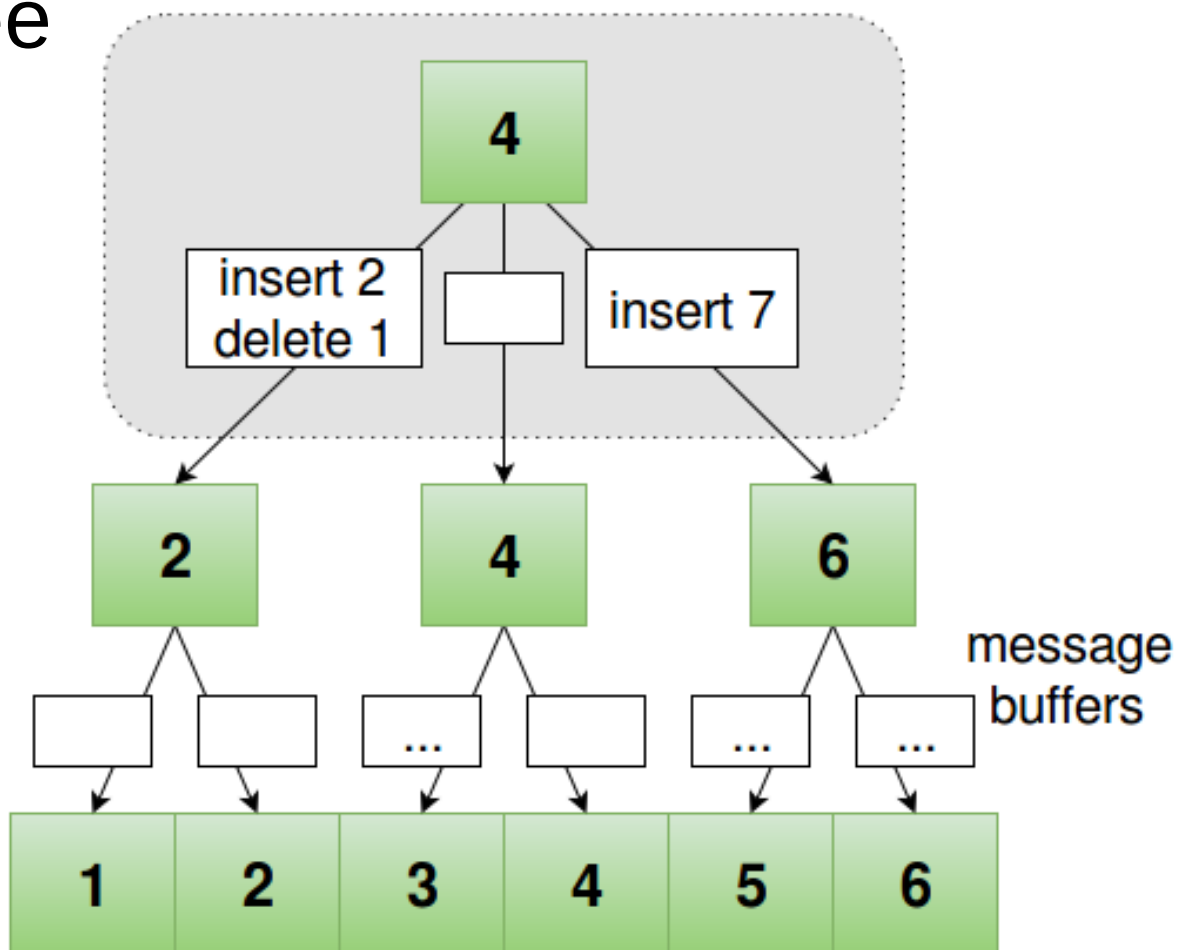
- Prototype shows incredible results!



- VACUUM?
- WAL?
- Concurrency?
- Index size?
- too hard =(

# Fractal Tree


- Insert the message instead of the data
- Send it down the tree
- Apply a message to leaf page



- TokuDB for MySQL
- TokuMX for MongoDB



# PostgreSQL specific



## PostgreSQL specific (1)


- Write-Ahead-Log
  - WAL is not extendable
- File manager
  - 1 Relation (Heap or Index) = 1 continuous file
  - Free Space Map
- Block size
  - 8 Kb

# Advanced indexes




- Optimize the number of indexes
  - `pg_stat_statements`
- REINDEX
- CREATE INDEX CONCURRENTLY
  - rebuild bloated and fragmented indexes
- Partial indexes
- BRIN
  - tiny min/max index

# Ideas?



# Covering and Unique


- To maintain constraint (Unique, Primary key..) on A
- To use IndexOnlyScan on A,B
- Have to maintain 2 indexes



# Covering + Unique


```
CREATE UNIQUE INDEX ON mytable
USING btree(a)
INCLUDING(b);
```

- DONE

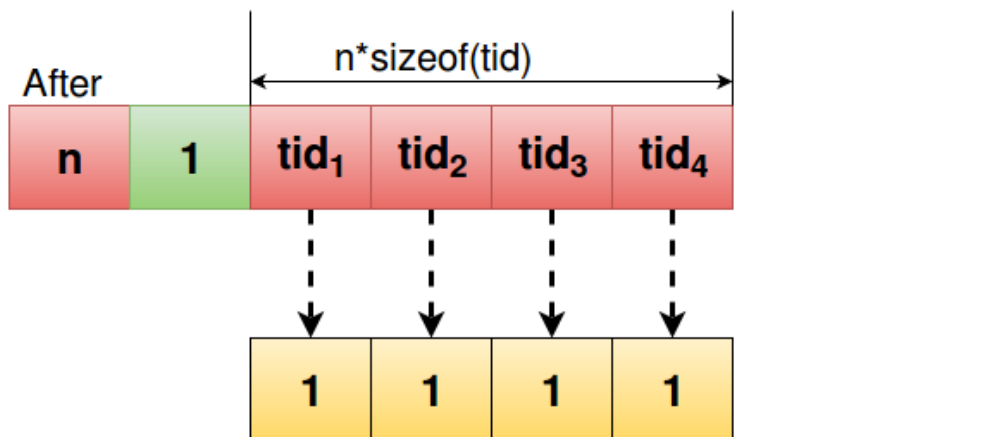


# Effective storage of duplicates

- Compress duplicated keys on index page



- IN PROGRESS

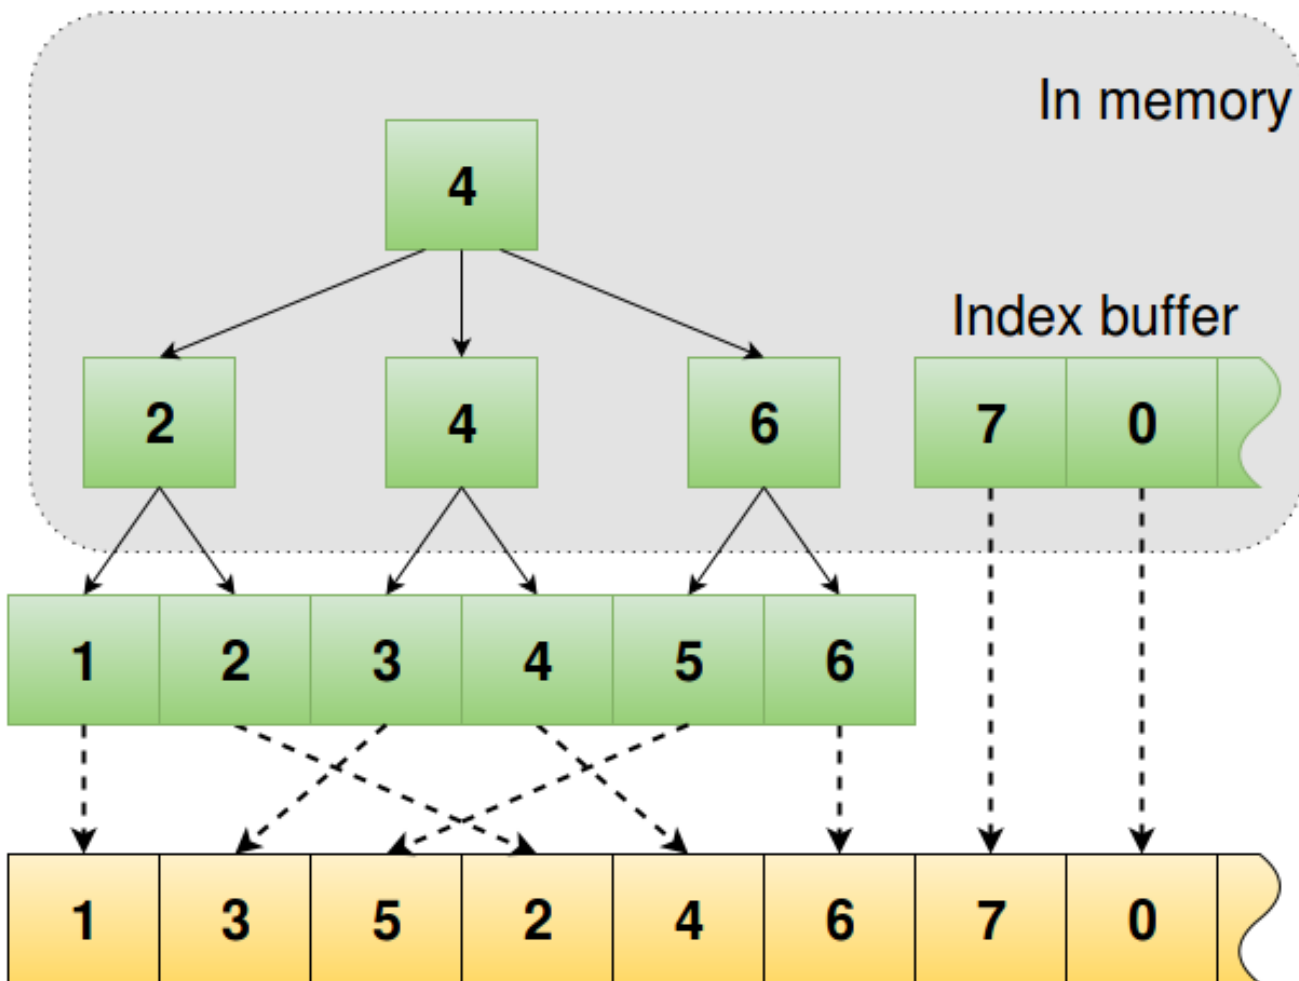




- INSERT INTO mytable  
SELECT x  
FROM generate\_series(0, 1000000) as x;
- 1.000.000 B-tree searches
- 1.000.000 WAL records
- TODO

# Insert Buffer

- Flexible
- Recoverable
- TODO





Thanks for attention!  
Any questions?

[a.lubennikova@postgrespro.ru](mailto:a.lubennikova@postgrespro.ru)

[www.postgrespro.ru](http://www.postgrespro.ru)